

COMP9024: Data Structures and Algorithms

Week Seven: Priority Queues

Hui Wu

Session 1, 2015

<http://www.cse.unsw.edu.au/~cs9024>

1

Outline

- Priority Queues
- Heaps
- Adaptable Priority Queues

2

Priority Queues



3

Priority Queue ADT

- A priority queue stores a collection of entries.
- Each **entry** is a pair (key, value).
- Main methods of the Priority Queue ADT:
 - `insert(k, x)`
Inserts an entry with key *k* and value *x*.
 - `removeMin()`
Removes and returns the entry with smallest key.
- Additional methods
 - `min()`
returns, but does not remove, an entry with smallest key
 - `size()`, `isEmpty()`
- Applications:
 - Standby flyers
 - Auctions
 - Stock market

4

Total Order Relations

- Keys in a priority queue can be arbitrary objects on which an order is defined.
- Two distinct entries in a priority queue can have the same key.
- Mathematical concept of total order relation \leq
 - Reflexive property:
 $x \leq x$
 - Antisymmetric property:
 $x \leq y \wedge y \leq x \Rightarrow x = y$
 - Transitive property:
 $x \leq y \wedge y \leq z \Rightarrow x \leq z$

5

Entry ADT

- An **entry** in a priority queue is simply a key-value pair.
- Priority queues store entries to allow for efficient insertion and removal based on keys.
- Methods:
 - `key()`: returns the key for this entry.
 - `value()`: returns the value associated with this entry.
- As a Java interface:

```
/**
 * Interface for a key-value
 * pair entry
 */
public interface Entry {
    public Object key();
    public Object value();
}
```

6

Comparator ADT

- A comparator encapsulates the action of comparing two objects according to a given total order relation.
- A generic priority queue uses an auxiliary comparator.
- The comparator is external to the keys being compared
- When the priority queue needs to compare two keys, it uses its comparator.
- The primary method of the Comparator ADT:
 - `compare(a, b)`: Returns an integer i such that $i < 0$ if $a < b$, $i = 0$ if $a = b$, and $i > 0$ if $a > b$; an error occurs if a and b cannot be compared.

7

Example Comparator

Lexicographic comparison of 2-D points:

```

/** Comparator for 2D points under the
    standard lexicographic order. */
public class Lexicographic implements
    Comparator {
    int xa, ya, xb, yb;
    public int compare(Object a, Object b)
        throws ClassCastException {
        xa = ((Point2D) a).getX();
        ya = ((Point2D) a).getY();
        xb = ((Point2D) b).getX();
        yb = ((Point2D) b).getY();
        if (xa != xb)
            return (xb - xa);
        else
            return (yb - ya);
    }
}

```

Point objects:

```

/** Class representing a point in the
    plane with integer coordinates */
public class Point2D {
    protected int xc, yc; // coordinates
    public Point2D(int x, int y) {
        xc = x;
        yc = y;
    }
    public int getX() {
        return xc;
    }
    public int getY() {
        return yc;
    }
}

```

8

Priority Queue Sorting

- We can use a priority queue to sort a set of comparable elements:
 - Insert the elements one by one with a series of `insert` operations.
 - Remove the elements in sorted order with a series of `removeMin` operations.
- The running time of this sorting method depends on the priority queue implementation.

Algorithm PQ-Sort(S, C)
Input sequence S , comparator C for the elements of S
Output sequence S sorted in increasing order according to C

```

{  $P$  = priority queue with comparator  $C$ ;
  while ( $\neg S.isEmpty()$ )
  {  $e = S.removeFirst()$ ;
     $P.insert(e, 0)$ ; }
  while ( $\neg P.isEmpty()$ )
  {  $e = P.removeMin().key()$ ;
     $S.insertLast(e)$ ; }
}

```

9

List-based Priority Queue

- Implementation with an unsorted list:

4 5 2 3 1

- Performance:
 - `insert` takes $O(1)$ time since we can insert the item at the beginning or end of the list.
 - `removeMin` and `min` take $O(n)$ time since we have to traverse the entire list to find the smallest key.

- Implementation with a sorted list:

1 2 3 4 5

- Performance:
 - `insert` takes $O(n)$ time since we have to find the place where to insert the item.
 - `removeMin` and `min` take $O(1)$ time, since the smallest key is at the beginning.

10

Selection-Sort

- Selection-sort is the variation of PQ-sort where the priority queue is implemented with an unsorted list.
- Running time of Selection-sort:
 - Inserting the elements into the priority queue with n `insert` operations takes $O(n)$ time.
 - Removing the elements in sorted order from the priority queue with n `removeMin` operations takes time proportional to $1 + 2 + \dots + n$
- Selection-sort runs in $O(n^2)$ time.

11

Selection-Sort Example

| | List S | Priority Queue P |
|---------|----------------------------|--------------------------------------|
| Input: | (7,4,8,2,5,3,9) | () |
| Phase 1 | | |
| (a) | (4,8,2,5,3,9) | (7) |
| (b) | (8,2,5,3,9) | (7,4) |
| .. | | |
| (g) | () | (7,4,8,2,5,3,9) |
| Phase 2 | | |
| (a) | (2) | (7,4,8,5,3,9) |
| (b) | (2,3) | (7,4,8,5,9) |
| (c) | (2,3,4) | (7,8,5,9) |
| (d) | (2,3,4,5) | (7,8,9) |
| (e) | (2,3,4,5,7) | (8,9) |
| (f) | (2,3,4,5,7,8) | (9) |
| (g) | (2,3,4,5,7,8,9) | () |

12

Insertion-Sort

- Insertion-sort is the variation of PQ-sort where the priority queue is implemented with a sorted list.
- Running time of Insertion-sort:
 - Inserting the elements into the priority queue with n `insert` operations takes time proportional to $1 + 2 + \dots + n$
 - Removing the elements in sorted order from the priority queue with a series of n `removeMin` operations takes $O(n)$ time.
- Insertion-sort runs in $O(n^2)$ time.

13


Insertion-Sort Example

| | List S | Priority queue P |
|---------|-----------------|--------------------|
| Input: | (7,4,8,2,5,3,9) | () |
| Phase 1 | | |
| (a) | (4,8,2,5,3,9) | (7) |
| (b) | (8,2,5,3,9) | (4,7) |
| (c) | (2,5,3,9) | (4,7,8) |
| (d) | (5,3,9) | (2,4,7,8) |
| (e) | (3,9) | (2,4,5,7,8) |
| (f) | (9) | (2,3,4,5,7,8) |
| (g) | () | (2,3,4,5,7,8,9) |
| Phase 2 | | |
| (a) | (2) | (3,4,5,7,8,9) |
| (b) | (2,3) | (4,5,7,8,9) |
| .. | .. | .. |
| (g) | (2,3,4,5,7,8,9) | () |

14


In-place Insertion-sort

- Instead of using an external data structure, we can implement selection-sort and insertion-sort in-place.
- A portion of the input list itself serves as the priority queue.
- For in-place insertion-sort
 - We keep sorted the initial portion of the list.
 - We can use `swaps` instead of modifying the list.



15

Heaps



16

Recall Priority Queue ADT

- A priority queue stores a collection of entries.
- Each **entry** is a pair (key, value).
- Main methods of the Priority Queue ADT:
 - `insert(k, x)` inserts an entry with key k and value x .
 - `removeMin()` removes and returns the entry with smallest key.
- Additional methods:
 - `min()` returns, but does not remove, an entry with smallest key
 - `size()`, `isEmpty()`
- Applications:
 - Standby flyers
 - Auctions
 - Stock market

17

Recall Priority Queue Sorting



- We can use a priority queue to sort a set of comparable elements:
 - Insert the elements with a series of `insert` operations
 - Remove the elements in sorted order with a series of `removeMin` operations
- The running time depends on the priority queue implementation:
 - Unsorted sequence gives selection-sort: $O(n^2)$ time
 - Sorted sequence gives insertion-sort: $O(n^2)$ time
- Can we do better?

Algorithm *PQ-Sort*(S, C)
Input sequence S , comparator C for the elements of S
Output sequence S sorted in increasing order according to C


```

{  $P$  = priority queue with comparator  $C$ ;
  while ( $\neg S.isEmpty()$ ) {
     $e = S.removeFirst()$ ;
     $P.insert(e, 0)$ ;
  }
  while ( $\neg P.isEmpty()$ ) {
     $e = P.removeMin().key()$ ;
     $S.insertLast(e)$ ;
  }
}
```

18

Heaps

- A heap is a binary tree storing keys at its nodes and satisfying the following properties:
 - Heap-Order:** for every internal node v other than the root, $key(v) \geq key(parent(v))$
 - Complete Binary Tree:** let h be the height of the heap
 - for $i = 0, \dots, h-1$, there are 2^i nodes of depth i
 - at depth $h-1$, the internal nodes are to the left of the external nodes
- The last node of a heap is the rightmost node of depth h .




19

Height of a Heap




- Theorem:** A heap storing n keys has height $O(\log n)$.
Proof: (we apply the complete binary tree property)
 - Let h be the height of a heap storing n keys.
 - Since there are 2^i keys at depth $i = 0, \dots, h-1$ and at least one key at depth h , we have $n \geq 1 + 2 + 4 + \dots + 2^{h-1} + 1$.
 - Thus, $n \geq 2^h$, i.e., $h \leq \log n$.



20

Heaps and Priority Queues


- We can use a heap to implement a priority queue.
- We store a (key, element) item at each internal node.
- We keep track of the position of the last node.
- For simplicity, we show only the keys in the pictures.



21

Insertion into a Heap


- Method insertItem of the priority queue ADT corresponds to the insertion of a key k to the heap.
- The insertion algorithm consists of three steps:
 - Find the insertion node z (the new last node)
 - Store k at z
 - Restore the heap-order property (discussed next)



22

Upheap


- After the insertion of a new key k , the heap-order property may be violated.
- Algorithm upheap restores the heap-order property by swapping k along an upward path from the insertion node.
- Upheap terminates when the key k reaches the root or a node whose parent has a key smaller than or equal to k .
- Since a heap has height $O(\log n)$, upheap runs in $O(\log n)$ time.



23

Removal from a Heap

- Method removeMin of the priority queue ADT corresponds to the removal of the root key from the heap.
- The removal algorithm consists of three steps:
 - Replace the root key with the key of the last node w
 - Remove w
 - Restore the heap-order property (discussed next)



24

Downheap


- After replacing the root key with the key k of the last node, the heap-order property may be violated.
- Algorithm downheap restores the heap-order property by swapping key k along a downward path from the root.
- Downheap terminates when key k reaches a leaf or a node whose children have keys greater than or equal to k .
- Since a heap has height $O(\log n)$, downheap runs in $O(\log n)$ time.



25

Updating the Last Node

- The insertion node can be found by traversing a path of $O(\log n)$ nodes:
 - Go up until a left child or the root is reached
 - If a left child is reached, go to the right child
 - Go down left until a leaf is reached
- Similar algorithm for updating the last node after a removal.



26

Heap-Sort




- Consider a priority queue with n items implemented by means of a heap
 - the space used is $O(n)$
 - methods `insert` and `removeMin` take $O(\log n)$ time
 - methods `size`, `isEmpty`, and `min` take time $O(1)$ time
- Using a heap-based priority queue, we can sort a sequence of n elements in $O(n \log n)$ time.
- The resulting algorithm is called heap-sort.
- Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort.

27

Array List-based Heap Implementation


- We can represent a heap with n keys by means of an array list of length $n + 1$.
- For the node at rank i
 - the left child is at rank $2i$
 - the right child is at rank $2i + 1$
- Links between nodes are not explicitly stored.
- The cell of at rank 0 is not used..
- Operation `insert` corresponds to inserting at rank $n + 1$.
- Operation `removeMin` corresponds to removing at rank n .
- Yields in-place heap-sort.



28

Merging Two Heaps

- We are given two two heaps and a key k .
- We create a new heap with the root node storing k and with the two heaps as subtrees
- We perform downheap to restore the heap-order property.




29

Bottom-up Heap Construction




- We can construct a heap storing n given keys in using a bottom-up construction with $\log n$ phases.
- In phase i , pairs of heaps with $2^i - 1$ keys are merged into heaps with $2^{i+1} - 1$ keys.




30

Example (1/4)




31

Example (2/4)




32

Example (3/4)



33


Example (4/4)



34


Analysis

- We visualize the worst-case time of a downheap with a proxy path that goes first right and then repeatedly goes left until the bottom of the heap (this path may differ from the actual downheap path)
- Since each node is traversed by at most two proxy paths, the total number of nodes of the proxy paths is $O(n)$.
- Thus, bottom-up heap construction runs in $O(n)$ time.
- Bottom-up heap construction is faster than n successive insertions and speeds up the first phase of heap-sort.



35

Adaptable Priority Queues



36

Recall the Entry and Priority Queue ADTs

- An **entry** stores a (key, value) pair within a data structure
- Methods of the entry ADT:
 - `key()`: returns the key associated with this entry
 - `value()`: returns the value paired with the key associated with this entry
- Priority Queue ADT:
 - `insert(k, x)` inserts an entry with key k and value x
 - `removeMin()` removes and returns the entry with smallest key
 - `min()` returns, but does not remove, an entry with smallest key
 - `size()`, `isEmpty()`

37

Motivating Example



- Suppose we have an online trading system where orders to purchase and sell a given stock are stored in two priority queues (one for sell orders and one for buy orders) as (p, s) entries:
 - The key, p , of an order is the price
 - The value, s , for an entry is the number of shares
 - A buy order (p, s) is executed when a sell order (p', s') with price $p' \leq p$ is added (the execution is complete if $s' \geq s$)
 - A sell order (p, s) is executed when a buy order (p', s') with price $p' \geq p$ is added (the execution is complete if $s' \geq s$)
- What if someone wishes to cancel their order before it executes?
- What if someone wishes to update the price or number of shares for their order?

38

Methods of the Adaptable Priority Queue ADT

- `remove(e)`: Remove from P and return entry e .
- `replaceKey(e, k)`: Replace with k and return the key of entry e of P ; an error condition occurs if k is invalid (that is, k cannot be compared with other keys).
- `replaceValue(e, x)`: Replace with x and return the value of entry e of P .

39

Example

| Operation | Output | P |
|-----------------------------------|--------|--------------------------|
| <code>insert(5, A)</code> | e_1 | $(5, A)$ |
| <code>insert(3, B)</code> | e_2 | $(3, B), (5, A)$ |
| <code>insert(7, C)</code> | e_3 | $(3, B), (5, A), (7, C)$ |
| <code>min()</code> | e_2 | $(3, B), (5, A), (7, C)$ |
| <code>key(e_2)</code> | 3 | $(3, B), (5, A), (7, C)$ |
| <code>remove(e_1)</code> | e_1 | $(3, B), (7, C)$ |
| <code>replaceKey(e_2, 9)</code> | 3 | $(7, C), (9, B)$ |
| <code>replaceValue(e_3, D)</code> | C | $(7, D), (9, B)$ |
| <code>remove(e_2)</code> | e_2 | $(7, D)$ |

40

Locating Entries

- In order to implement the operations `remove(k)`, `replaceKey(e)`, and `replaceValue(k)`, we need fast ways of locating an entry e in a priority queue.
- We can always just search the entire data structure to find an entry e , but there are better ways for locating entries.

41

Location-Aware Entries




- A locator-aware entry identifies and tracks the location of its (key, value) object within a data structure
- Intuitive notion:
 - Coat claim check
 - Valet claim ticket
 - Reservation number
- Main idea:
 - Since entries are created and returned from the data structure itself, it can return location-aware entries, thereby making future updates easier

42

List Implementation


- A location-aware list entry is an object storing
 - key
 - value
 - position (or rank) of the item in the list
- In turn, the position (or array cell) stores the entry
- Back pointers (or ranks) are updated during swaps



43

Heap Implementation

- A location-aware heap entry is an object storing
 - key
 - value
 - position of the entry in the underlying heap
- In turn, each heap position stores an entry
- Back pointers are updated during entry swaps



44

Performance

- Using location-aware entries we can achieve the following running times (times better than those achievable without location-aware entries are highlighted in blue):

| Method | Unsorted List | Sorted List | Heap |
|---------------|---------------|-------------|-------------|
| size, isEmpty | $O(1)$ | $O(1)$ | $O(1)$ |
| insert | $O(1)$ | $O(n)$ | $O(\log n)$ |
| min | $O(n)$ | $O(1)$ | $O(1)$ |
| removeMin | $O(n)$ | $O(1)$ | $O(\log n)$ |
| remove | $O(1)$ | $O(1)$ | $O(\log n)$ |
| replaceKey | $O(1)$ | $O(n)$ | $O(\log n)$ |
| replaceValue | $O(1)$ | $O(1)$ | $O(1)$ |

45

References

1. Chapter 8, Data Structures and Algorithms by Goodrich and Tamassia.

46