

COMP9024: Data Structures and Algorithms

Week Ten: Text Processing

Hui Wu

Session 1, 2015

<http://www.cse.unsw.edu.au/~cs9024>

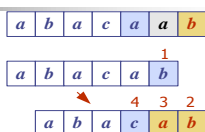
1

Outline

- Pattern Matching
- Tries
- The Greedy Method and Text Compression
- Dynamic Programming

2

Pattern Matching



3

Strings

- A string is a sequence of characters
- Examples of strings:
 - Java program
 - HTML document
 - DNA sequence
 - Digitized image
- An alphabet Σ is the set of possible characters for a family of strings
- Example of alphabets:
 - ASCII
 - Unicode
 - $\{0, 1\}$
 - $\{A, C, G, T\}$
- Let P be a string of size m
 - A substring $P[i..j]$ of P is the subsequence of P consisting of the characters with ranks between i and j
 - A prefix of P is a substring of the type $P[0..i]$
 - A suffix of P is a substring of the type $P[i..m-1]$
- Given strings T (text) and P (pattern), the pattern matching problem consists of finding a substring of T equal to P
- Applications:
 - Text editors
 - Search engines
 - Biological research

4

Brute-Force Pattern Matching

- The brute-force pattern matching algorithm compares the pattern P with the text T for each possible shift of P relative to T , until either
 - a match is found, or
 - all placements of the pattern have been tried
- Brute-force pattern matching runs in time $O(nm)$
- Example of worst case:
 - $T = aaa \dots ah$
 - $P = aaah$
 - may occur in images and DNA sequences
 - unlikely in English text

Algorithm BruteForceMatch(T, P)
Input text T of size n and pattern P of size m
Output starting index of a substring of T equal to P or -1 if no such substring exists

```
{ for ( $i = 0$ ;  $i < n - m + 1$ ;  $i++$ )  
  { // test shift  $i$  of the pattern  
     $j = 0$ ;  
    while ( $j < m \wedge T[i + j] = P[j]$ )  
       $j = j + 1$ ;  
    if ( $j = m$ )  
      return  $i$ ; // match at  $i$   
  }  
return  $-1$  // no match anywhere  
}
```

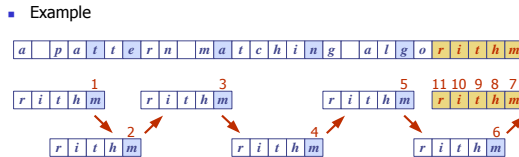
5

Boyer-Moore Heuristics

The Boyer-Moore's pattern matching algorithm is based on two heuristics

Looking-glass heuristic: Compare P with a subsequence of T moving backwards

- **Character-jump heuristic:** When a mismatch occurs at $T[i] = c$
 - If P contains c , shift P to align the last occurrence of c in P with $T[i]$
 - Else, shift P to align $P[0]$ with $T[i + 1]$



6

Last-Occurrence Function

- Boyer-Moore's algorithm preprocesses the pattern P and the alphabet Σ to build the last-occurrence function L mapping Σ to integers, where $L(c)$ is defined as
 - the largest index i such that $P[i] = c$ or
 - 1 if no such index exists
- Example:

c	a	b	c	d
$L(c)$	4	5	3	-1

 - $\Sigma = \{a, b, c, d\}$
 - $P = abacab$
- The last-occurrence function can be represented by an array indexed by the numeric codes of the characters
- The last-occurrence function can be computed in time $O(m + s)$, where m is the size of P and s is the size of Σ

c	a	b	c	d
$L(c)$	4	5	3	-1

7

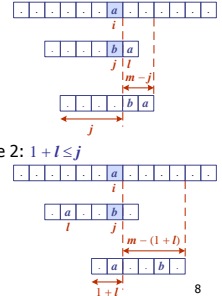
The Boyer-Moore Algorithm

```

Algorithm BoyerMooreMatch( $T, P, \Sigma$ )
{  $L = \text{lastOccurrenceFunction}(P, \Sigma)$ 
   $i = m - 1$ 
   $j = m - 1$ 
  repeat
    if (  $T[i] = P[j]$  )
      { if (  $j = 0$  )
          return  $i$  // match at  $i$ 
        else
          {  $i = i - 1$ ;
             $j = j - 1$ ; }
      }
    else // character-jump
      {  $i = L[T[i]]$ ;
         $i = i + m - \min(j, 1 + i)$ ;
         $j = m - 1$ ; }
  until (  $i > n - 1$  )
  return  $-1$  // no match
}

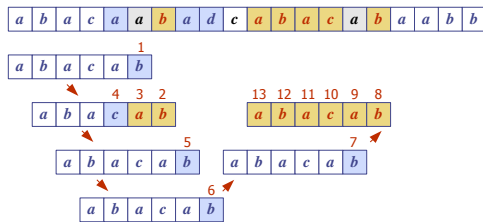
```

~~Case 1: $j \leq 1 + l$~~



8


Example



9

Analysis

- Boyer-Moore's algorithm runs in time $O(nm + s)$
- Example of worst case:
 - $T = aaa \dots a$
 - $P = baaa$
- The worst case may occur in images and DNA sequences but is unlikely in English text
- Boyer-Moore's algorithm is significantly faster than the brute-force algorithm on English text

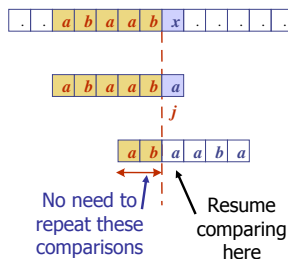


n
t

10

The KMP Algorithm

- Knuth-Morris-Pratt's algorithm compares the pattern to the text in **left-to-right**, but shifts the pattern more intelligently than the brute-force algorithm.
- When a mismatch occurs, what is the **most** we can shift the pattern so as to avoid redundant comparisons?
- Answer: the largest prefix of $P[0..j]$ that is a suffix of $P[1..j]$

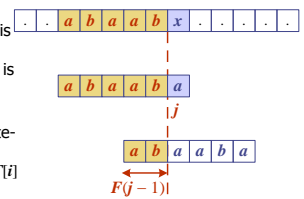


11

KMP Failure Function

- Knuth-Morris-Pratt's algorithm preprocesses the pattern to find matches of prefixes of the pattern with the pattern itself
- The **failure function** $F(j)$ is defined as the size of the largest prefix of $P[0..j]$ that is also a suffix of $P[1..j]$
- Knuth-Morris-Pratt's algorithm modifies the brute-force algorithm so that if a mismatch occurs at $P[j] \neq T[i]$ we set $j \leftarrow F(j - 1)$

j	0	1	2	3	4	5
$P[j]$	a	b	a	a	b	a
$F(j)$	0	0	1	1	2	3



12

The KMP Algorithm

- The failure function can be represented by an array and can be computed in $O(m)$ time
- At each iteration of the while-loop, either
 - i increases by one, or
 - the shift amount $i - j$ increases by at least one (observe that $F[j - 1] < j$)
- Hence, there are no more than $2n$ iterations of the while-loop
- Thus, KMP's algorithm runs in optimal time $O(m + n)$

```

Algorithm KMPMatch(P, T)
{
    F = failureFunction(P);
    i = 0;
    j = 0;
    while (i < n)
        if (T[i] = P[j])
            if (j = m - 1)
                return i - j; // match
            else
                { i = i + 1; j = j + 1; }
        else
            if (j > 0)
                j = F[j - 1];
            else
                i = i + 1;
    return -1; // no match
}
    
```

13

Computing the Failure Function



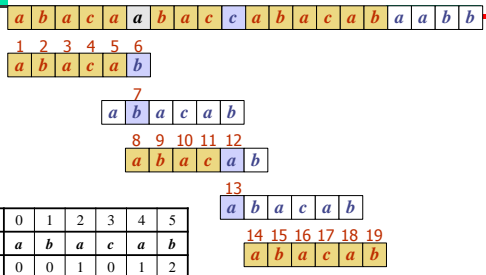
- The failure function can be represented by an array and can be computed in $O(m)$ time
- The construction is similar to the KMP algorithm itself
- At each iteration of the while-loop, either
 - i increases by one, or
 - the shift amount $i - j$ increases by at least one (observe that $F[j - 1] < j$)
- Hence, there are no more than $2m$ iterations of the while-loop

```

Algorithm failureFunction(P)
{
    F[0] = 0;
    i = 1;
    j = 0;
    while (i < m)
        if (P[i] = P[j])
            { // we have matched j + 1 char
                F[i] = j + 1;
                i = i + 1;
                j = j + 1; }
        else if (j > 0)
            // use failure function to shift P
            j = F[j - 1];
        else
            { F[i] = 0; // no match
              i = i + 1; }
    }
    
```

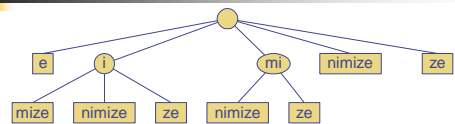
14

Example



15

Tries



16

Preprocessing Strings

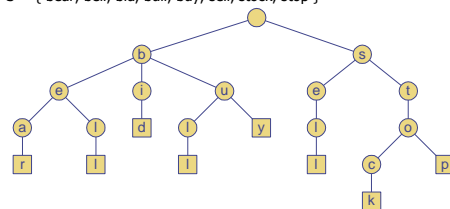
- Preprocessing the pattern speeds up pattern matching queries
 - After preprocessing the pattern, KMP's algorithm performs pattern matching in time proportional to the text size
- If the text is large, immutable and searched for often (e.g., works by Shakespeare), we may want to preprocess the text instead of the pattern
- A trie is a compact data structure for representing a set of strings, such as all the words in a text
 - A trie supports pattern matching queries in time proportional to the pattern size

17

Standard Tries

The standard trie for a set of strings S is an ordered tree such that:

- Each node but the root is labeled with a character
 - The children of a node are alphabetically ordered
 - The paths from the external nodes to the root yield the strings of S
- Example: standard trie for the set of strings
 $S = \{ \text{bear, bell, bid, bull, buy, sell, stock, stop} \}$

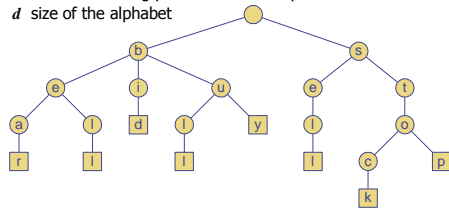


18

Analysis of Standard Tries

A standard trie uses $O(n)$ space and supports searches, insertions and deletions in time $O(dm)$, where:

- n total size of the strings in S
- m size of the string parameter of the operation
- d size of the alphabet

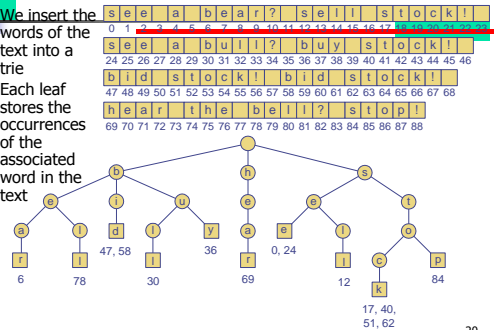


19

Word Matching with a Trie

We insert the words of the text into a trie

- Each leaf stores the occurrences of the associated word in the text

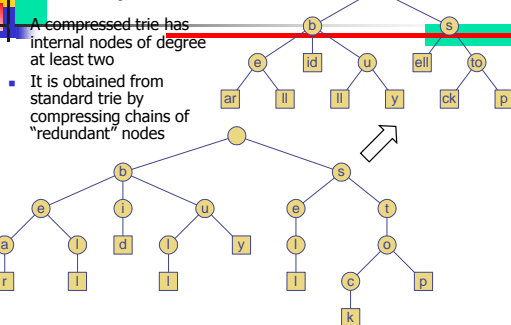


20

Compressed Tries

A compressed trie has internal nodes of degree at least two

- It is obtained from standard trie by compressing chains of "redundant" nodes

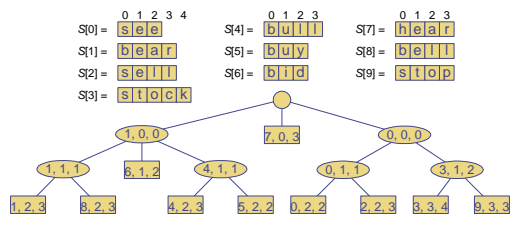


21

Compact Representation

Compact representation of a compressed trie for an array of strings:

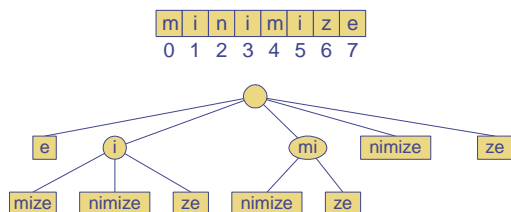
- Stores at the nodes ranges of indices instead of substrings
- Uses $O(s)$ space, where s is the number of strings in the array
- Serves as an auxiliary index structure



22

Suffix Trie

- The suffix trie of a string X is the compressed trie of all the suffixes of X



23

Pattern Matching Using Suffix Trie (1/2)

```

Algorithm suffixTrieMatch(w, k)
{
  p = P.Length; j = 0; v = T.Root();
  repeat
  {
    f = true;
    for each child w of v do
    {
      // we have matched j + 1 char
      i = start(w); // start(w) is the start index of w
      if ( P[j] = T[i] ) // process child w
      {
        x = end(w) - i + 1; // end(w) is the end index of w
        if ( p ≤ x )
          // suffix is shorter than or of the same length of the node label
          {
            if ( P[j+p-x-1] = X[i+p-x-1] ) return i - j;
            else return "P is not a substring of X";
          }
        else // suffix is longer than the node label
          {
            if ( P[j+p-x-1] = X[i+p-x-1] )
            {
              p = p - x; // update suffix length
              j = j + x; // update suffix start index
              v = w; f = false; break out of the for loop; } } }
            until f or T.isExternal(v);
            return "P is not a substring of X";
          }
    }
  }

```

24

Pattern Matching Using Suffix Trie (2/2)

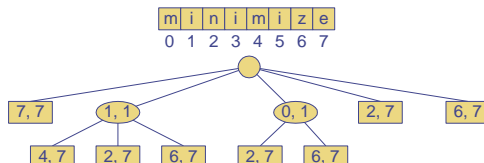
- Input of the algorithm:
 - Compact suffix trie T for a text X and pattern P .
- Output of the algorithm:
 - Starting index of a substring of X matching P or an indication that P is not a substring.
- The algorithm assumes the following additional property on the labels of the nodes in the compact representation of the suffix trie:
 - If node v has label (i, j) and Y is the string of length j associated with the path from the root to v (included), then $X[j-y+1..j]=Y$.
- This property ensures that we can easily compute the start index of the pattern in the text when a match occurs.

25

Analysis of Suffix Tries

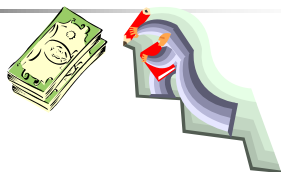
Compact representation of the suffix trie for a string X of size n from an alphabet of size d

- Uses $O(n)$ space
- Supports arbitrary pattern matching queries in X in $O(dm)$ time, where m is the size of the pattern
- Can be constructed in $O(n)$ time



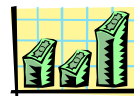
26

The Greedy Method and Text Compression



27

The Greedy Method Technique



The **greedy method** is a general algorithm design paradigm, built on the following elements:

- configurations**: different choices, collections, or values to find
- objective function**: a score assigned to configurations, which we want to either maximize or minimize
- It works best when applied to problems with the **greedy-choice** property:
 - a globally-optimal solution can always be found by a series of local improvements from a starting configuration.

28

Text Compression

- Given a string X , efficiently encode X into a smaller string Y
 - Saves memory and/or bandwidth
- A good approach: **Huffman encoding**
 - Compute frequency $f(c)$ for each character c .
 - Encode high-frequency characters with short code words
 - No code word is a prefix for another code
 - Use an optimal encoding tree to determine the code words

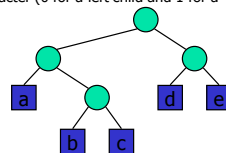
29

Encoding Tree Example

A **code** is a mapping of each character of an alphabet to a binary code-word

- A **prefix code** is a binary code such that no code-word is the prefix of another code-word
- An **encoding tree** represents a prefix code
 - Each external node stores a character
 - The code word of a character is given by the path from the root to the external node storing the character (0 for a left child and 1 for a right child)

00	010	011	10	11
a	b	c	d	e



30

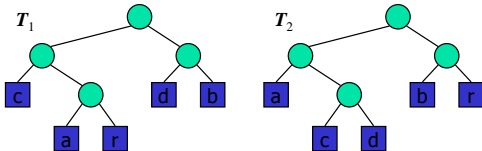
Encoding Tree Optimization

Given a text string X , we want to find a prefix code for the characters of X that yields a small encoding for X

- Frequent characters should have short code-words
- Rare characters should have long code-words

Example

- $X = \text{abracadabra}$
- T_1 encodes X into 29 bits
- T_2 encodes X into 24 bits



31

Huffman's Algorithm

Given a string X , Huffman's algorithm constructs a prefix code that minimizes the size of the encoding of X

- It runs in time $O(n + d \log d)$, where n is the size of X and d is the number of distinct characters of X
- A heap-based priority queue is used as an auxiliary structure

Algorithm *HuffmanEncoding*(X)

```

Input string  $X$  of size  $n$ 
Output optimal encoding tree for  $X$ 
{
   $C = \text{distinctCharacters}(X)$ ;
   $\text{computeFrequencies}(C, X)$ ;
   $Q = \text{new empty heap}$ ;
  for all  $c \in C$ 
  {  $T = \text{new single-node tree storing } c$ ;
     $Q.\text{insert}(\text{getFrequency}(c), T)$ ; }
  while ( $Q.\text{size}() > 1$ )
  {  $f_1 = Q.\text{minKey}()$ ;
     $T_1 = Q.\text{removeMin}()$ ;
     $f_2 = Q.\text{minKey}()$ ;
     $T_2 = Q.\text{removeMin}()$ ;
     $T = \text{join}(T_1, T_2)$ ;
     $Q.\text{insert}(f_1 + f_2, T)$ ;
  }
  return  $Q.\text{removeMin}()$ ;
}

```

32

Example

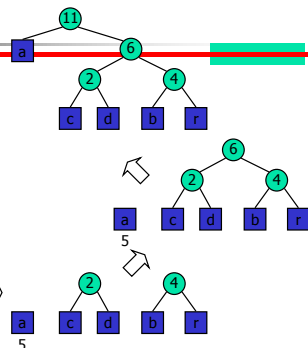
$X = \text{abracadabra}$

Frequencies

a	b	c	d	r
5	2	1	1	2

a	b	c	d	r
5	2	1	1	2

a	b	c	d	r
5	2	1	1	2

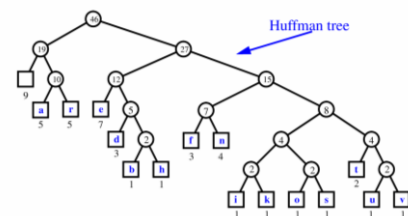


33

Extended Huffman Tree Example

String: a fast runner need never be afraid of the dark

Character	a	b	d	e	f	h	i	k	n	o	r	s	t	u	v	
Frequency	9	5	1	3	7	3	1	1	1	4	1	5	1	2	1	1



34

The Fractional Knapsack Problem (not in book)



- Given: A set S of n items, with each item i having
 - b_i - a positive benefit
 - w_i - a positive weight
- Goal: Choose items with maximum total benefit but with weight at most W .
- If we are allowed to take fractional amounts, then this is the **fractional knapsack problem**.
 - In this case, we let x_i denote the amount we take of item i

Objective: maximize $\sum_{i \in S} b_i(x_i / w_i)$

Constraint: $\sum_{i \in S} x_i \leq W$

35

Example



- Given: A set S of n items, with each item i having
 - b_i - a positive benefit
 - w_i - a positive weight
- Goal: Choose items with maximum total benefit but with weight at most W .

Items:	1	2	3	4	5
Weight:	4 ml	8 ml	2 ml	6 ml	1 ml
Benefit:	\$12	\$32	\$40	\$30	\$50
Value: (\$ per ml)	3	4	20	5	50



"knapsack"
Solution:
• 1 ml of 5
• 2 ml of 3
• 6 ml of 4
• 1 ml of 2

36

The Fractional Knapsack Algorithm



Greedy choice: **Keep taking** item with highest **value** (benefit to weight ratio)

- Since $\sum_{i \in S} b_i(x_i / w_i) = \sum_{i \in S} (b_i / w_i) x_i$
- Run time: $O(n \log n)$. Why?

Correctness: Suppose there is a better solution

- there is an item i with higher value than a chosen item j , but $x_i < w_i$, $x_j > 0$ and $v_i < v_j$
- If we substitute some i with j , we get a better solution
- How much of i : $\min\{w_i - x_i, x_j\}$
- Thus, there is no better solution than the greedy one

Algorithm fractionalKnapsack(S, W)
Input: set S of items with benefit b_i and weight w_i ; max. weight W
Output: amount x_i of each item i to maximize benefit with weight at most W

```

{ for each item  $i$  in  $S$ 
  {  $x_i = 0$ ;
     $v_i = b_i / w_i$ ; // value
  }
 $w = 0$ ; // total weight
while ( $w < W$ )
{ remove item  $i$  with highest  $v_i$ 
   $x_i = \min\{w_i, W - w\}$ ;
   $w = w + \min\{w_i, W - w\}$ ;
}

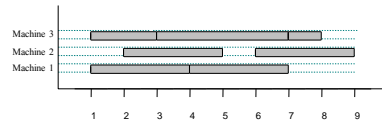
```

37

Task Scheduling (not in book)



- Given: a set T of n tasks, each having:
 - A start time, s_i
 - A finish time, f_i (where $s_i < f_i$)
- Goal: Perform all the tasks using a minimum number of "machines."



38

Task Scheduling Algorithm



Greedy choice: consider tasks by their start time and use as few machines as possible with this order.

- Run time: $O(n \log n)$. Why?
- Correctness: Suppose there is a better schedule.
 - We can use $k-1$ machines
 - The algorithm uses k
 - Let i be first task scheduled on machine k
 - Machine i must conflict with $k-1$ other tasks
 - But that means there is no non-conflicting schedule using $k-1$ machines

Algorithm taskSchedule(T)
Input: set T of tasks with start time s_i and finish time f_i
Output: non-conflicting schedule with minimum number of machines

```

{  $m = 0$ ; // no. of machines
  while  $T$  is not empty
  { remove task  $i$  with smallest  $s_i$ 
    if there's a machine  $j$  for  $i$  then
      schedule  $i$  on machine  $j$ ;
    else
      {  $m = m + 1$ ;
        schedule  $i$  on machine  $m$ ;
      }
  }
}

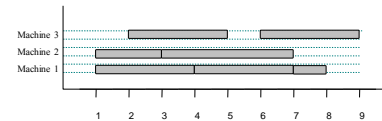
```

39

Example

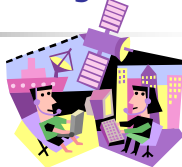


- Given: a set T of n tasks, each having:
 - A start time, s_i
 - A finish time, f_i (where $s_i < f_i$)
 - $[1,4], [1,3], [2,5], [3,7], [4,7], [6,9], [7,8]$ (ordered by start)
- Goal: Perform all tasks on min. number of machines



40

Dynamic Programming



41

Matrix Chain-Products (not in book)



- Dynamic Programming is a general algorithm design paradigm.

- Rather than give the general structure, let us first give a motivating example:

Matrix Chain-Products

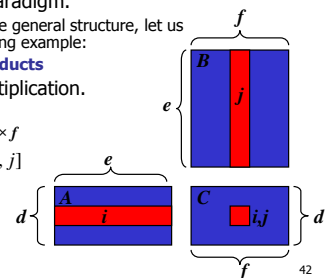
- Review: Matrix Multiplication.

- $C = A * B$

- A is $d \times e$ and B is $e \times f$

$$C[i, j] = \sum_{k=0}^{e-1} A[i, k] * B[k, j]$$

- $O(def)$ time



42

Matrix Chain-Products



Matrix Chain-Product:

- Compute $A = A_0 * A_1 * \dots * A_{n-1}$
- A_i is $d_i \times d_{i+1}$
- Problem: How to parenthesize?

Example

- B is 3×100
- C is 100×5
- D is 5×5
- $(B * C) * D$ takes $1500 + 75 = 1575$ ops
- $B * (C * D)$ takes $1500 + 2500 = 4000$ ops

43

An Enumeration Approach



Matrix Chain-Product Alg.:

- Try all possible ways to parenthesize
 $A = A_0 * A_1 * \dots * A_{n-1}$
- Calculate number of ops for each one
- Pick the one that is best
- Running time:
 - The number of parenthesizations is equal to the number of binary trees with n nodes
 - This is **exponential**!
 - It is called the Catalan number, and it is almost 4^n .
 - This is a terrible algorithm!

44

A Greedy Approach



- Idea #1: repeatedly select the product that uses (up) the most operations.
- Counter-example:
 - A is 10×5
 - B is 5×10
 - C is 10×5
 - D is 5×10
 - Greedy idea #1 gives $(A * B) * (C * D)$, which takes $500 + 1000 + 500 = 2000$ ops
 - $A * ((B * C) * D)$ takes $500 + 250 + 250 = 1000$ ops

45

Another Greedy Approach



- Idea #2: repeatedly select the product that uses the fewest operations.
- Counter-example:
 - A is 101×11
 - B is 11×9
 - C is 9×100
 - D is 100×99
 - Greedy idea #2 gives $A * ((B * C) * D)$, which takes $109989 + 9900 + 108900 = 228789$ ops
 - $(A * B) * (C * D)$ takes $9999 + 89991 + 89100 = 189090$ ops
- The greedy approach is not giving us the optimal value.

46

A "Recursive" Approach



- Define **subproblems**:
 - Find the best parenthesization of $A_i * A_{i+1} * \dots * A_j$.
 - Let $N_{i,j}$ denote the number of operations done by this subproblem.
 - The optimal solution for the whole problem is $N_{0,n-1}$.
- Subproblem optimality**: The optimal solution can be defined in terms of optimal subproblems
 - There has to be a final multiplication (root of the expression tree) for the optimal solution.
 - Say, the final multiply is at index i : $(A_0 * \dots * A_i) * (A_{i+1} * \dots * A_{n-1})$.
 - Then the optimal solution $N_{0,n-1}$ is the sum of two optimal subproblems, $N_{0,i}$ and $N_{i+1,n-1}$ plus the time for the last multiply.
 - If the global optimum did not have these optimal subproblems, we could define an even better "optimal" solution.

47

A Characterizing Equation



- The global optimal has to be defined in terms of optimal subproblems, depending on where the final multiply is at.
- Let us consider all possible places for that final multiply:
 - Recall that A_i is a $d_i \times d_{i+1}$ dimensional matrix.
 - So, a characterizing equation for $N_{i,j}$ is the following:

$$N_{i,j} = \min_{i \leq k < j} \{ N_{i,k} + N_{k+1,j} + d_i d_{k+1} d_{j+1} \}$$

- Note that subproblems are not independent--the **subproblems overlap**.

48

A Dynamic Programming Algorithm



- Since subproblems overlap, we don't use recursion.
- Instead, we construct optimal subproblems "bottom-up."
- N_{ij} 's are easy, so start with them
- Then do length 2, 3, ... subproblems, and so on.
- Running time: $O(n^3)$

Algorithm *matrixChain(S)*:

Input: sequence S of n matrices to be multiplied

Output: number of operations in an optimal parenthesization of S

```
{ for (i = 1; i ≤ n-1; i++)
  Nii = 0;
  for (b = 1; b ≤ n-i; b++)
    for (i = 0; i ≤ n-b-1; i++)
      { j = i+b;
        Nij = +infinity;
        for (k = i; k ≤ j-1; k++)
          Nij = min{Nij, Nik + Nk+1,j + didk+1
                    dj+1};
      }
```

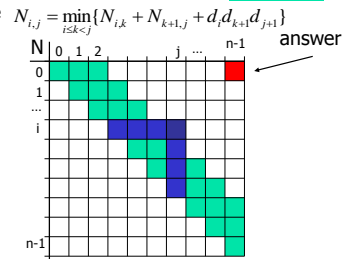
49

A Dynamic Programming Algorithm Visualization



The bottom-up construction fills in the N array by diagonals

- N_{ij} gets values from previous entries in i -th row and j -th column
- Filling in each entry in the N table takes $O(n)$ time.
- Total running time: $O(n^3)$
- Getting actual parenthesization can be done by remembering "k" for each N entry



50

The General Dynamic Programming Technique



- Applies to a problem that at first seems to require a lot of time (possibly exponential), provided we have:
 - Simple subproblems:** the subproblems can be defined in terms of a few variables, such as j , k , l , m , and so on.
 - Subproblem optimality:** the global optimum value can be defined in terms of optimal subproblems
 - Subproblem overlap:** the subproblems are not independent, but instead they overlap (hence, should be constructed bottom-up).

51

Subsequences

- A **subsequence** of a character string $x_0x_1x_2\dots x_{n-1}$ is a string of the form $x_{i_1}x_{i_2}\dots x_{i_k}$, where $i_j < i_{j+1}$.
- Not the same as substring!
- Example String: ABCDEFGHIJK
 - Subsequence: ACEGJIK
 - Subsequence: DFGHK
 - Not subsequence: DAGH

52

The Longest Common Subsequence (LCS) Problem

- Given two strings X and Y , the longest common subsequence (LCS) problem is to find a longest subsequence common to both X and Y
- Has applications to DNA similarity testing (alphabet is $\{A, C, G, T\}$)
- Example: ABCDEFG and XZACKDFWGH have ACDFG as a longest common subsequence

53

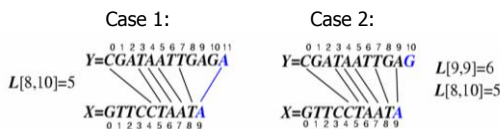
A Poor Approach to the LCS Problem

- A Brute-force solution:
 - Enumerate all subsequences of X
 - Test which ones are also subsequences of Y
 - Pick the longest one.
- Analysis:
 - If X is of length n , then it has 2^n subsequences
 - This is an exponential-time algorithm!

54

A Dynamic-Programming Approach to the LCS Problem

- Define $L[i,j]$ to be the length of the longest common subsequence of $X[0..i]$ and $Y[0..j]$.
- Allow for -1 as an index, so $L[-1,k] = 0$ and $L[k,-1] = 0$, to indicate that the null part of X or Y has no match with the other.
- Then we can define $L[i,j]$ in the general case as follows:
 - If $x_i = y_j$, then $L[i,j] = L[i-1,j-1] + 1$ (we can add this match)
 - If $x_i \neq y_j$, then $L[i,j] = \max\{L[i-1,j], L[i,j-1]\}$ (we have no match here)



An LCS Algorithm

Algorithm $LCS(X, Y)$:

Input: Strings X and Y with n and m elements, respectively

Output: For $i = 0, \dots, n-1$, $j = 0, \dots, m-1$, the length $L[i,j]$ of a longest string that is a subsequence of both the string $X[0..i] = x_0x_1x_2\dots x_i$ and the string $Y[0..j] = y_0y_1y_2\dots y_j$

```

{ for ( i=1; i ≤ n-1, i++ )
    L[i,-1] = 0;
  for ( j=0; j ≤ m-1, j++ )
    L[-1,j] = 0;
  for ( i=0; i ≤ n-1, i++ )
    for ( j=0; j ≤ m-1, j++ )
      if (  $x_i = y_j$  )
        L[i,j] = L[i-1,j-1] + 1;
      else
        L[i,j] = max{L[i-1,j], L[i,j-1]};
  return array L;
}
```

56

Visualizing the LCS Algorithm



57

Analysis of LCS Algorithm

- We have two nested loops
 - The outer one iterates n times
 - The inner one iterates m times
 - A constant amount of work is done inside each iteration of the inner loop
 - Thus, the total running time is $O(nm)$
- Answer is contained in $L[n,m]$ (and the subsequence can be recovered from the L table).

58

References

- Chapter 12, Data Structures and Algorithms by Goodrich and Tamassia.

59